

SANJANA VENKATESH

Los Angeles, CA | (213) 913-1568 | sv51119@usc.edu | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EDUCATION

University of Southern California Master of Science in Analytics GPA: 3.8/4.00 Relevant Coursework: Data Analysis, Software Engineering, Data Warehousing & ETL/ELT, Cloud Analytics	Los Angeles, CA Aug 2023-May 2025
BMS College of Engineering Bachelor of Technology in Computer Science; GPA: 3.7/4.00	Bengaluru, India Jul 2017 - Jul 2021

WORK EXPERIENCE

University of Southern California Data Analyst Research Assistant	Los Angeles, CA May 2025 – Present
<ul style="list-style-type: none">Orchestrated ML pipelines on Vertex AI and Snowflake, added CI/CD & Git to automate ops, cutting training time by 2hrsCrafted Tableau EDA and engineered features; optimized BigQuery (partitioning, clustering) to reduce query cost 38%	
SAG-AFTRA Health Plan Data Analyst Intern	Burbank, CA May 2024 – Aug 2024
<ul style="list-style-type: none">Managed AWS resources for Big Data Analytics, optimising EC2 Instances, reducing operational cost by 23k annuallyAutomated ingestion pipelines from ThoughtSpot APIs into Snowflake using REST APIs, streamlining ETL operations and reducing latency by 42%	
Quantiphi, Inc Data Engineer	Bengaluru, India Feb 2021 – Jul 2023
<ul style="list-style-type: none">Modernized a legacy Hadoop estate on Google Cloud: migrated Hadoop/MapReduce jobs to Dataproc, tuned Spark (executor/EFM parameters, SSD/shuffle balancing), yielding \$10M/year infrastructure savingsDirected an 8-engineer program to migrate 1,500+ tables to BigQuery; hardened data quality with SQL checks and DVT, achieving ≥99% accuracy and audit-ready reporting complianceProductionized ETL/ELT pipelines with Apache Airflow from GCS to BigQuery, reducing manual ops by 16 hours/day and consistently meeting daily refresh SLAsOperationalized a governed metrics layer: authored LookML and launched energy KPI dashboards that drove \$1.2M annual cost savings; presented results to client VPs, enabling cross-functional adoptionImplemented a real-time pipeline to process 150M raw records from 30+ sources using Kafka and Pyspark	

PROJECTS

Recommendation Engine Link	Aug 2025 – Sep 2025
<ul style="list-style-type: none">Engineered a RAG recommender with LangChain, OpenAI embeddings, LanceDB to index titles/genres/synopses and return top-N semantic recs from NL queries and increased user engagement by 25%Packaged an end-to-end, reproducible pipeline to speed LLM-for-RecSys experiments and enable plug-and-play datasets	
Retail Data Analytics Link	Mar 2024 – Jul 2024
<ul style="list-style-type: none">Automated ETL with Apache Airflow + Python to ingest GCS to BigQuery, slashing data latency from 2 hours to 12 minutes and enabling reliable daily refresh SLAsBuilt the analytics layer: dbt star schema with data quality tests, partitioning & clustering for performance; delivered Power BI dashboards with DAX for Top Products /Markets, reducing manual reporting time by 48%	
Intelligent Fab Monitoring & Maintenance System Link	Jan 2024 – Mar 2024
<ul style="list-style-type: none">Built a real-time geospatial analytics platform tracking 2M+ devices ; delivered KPI dashboards, interactive data visualization, and event-driven alerts, eliminating manual monitoringEngineered ETL/ELT for 100K+ RTLS logs using Python performed data cleaning, geofencing, feature engineering and time-series modeling; trained XGBoost for RUL prediction (82% accuracy). 1st place among 20+ USC capstone teams	

TECHNICAL SKILLS

Languages: Python; Advanced SQL (BigQuery, Postgres); Bash; PySpark; **Certification:** Associate Cloud Engineer (Google Cloud)
Big Data & Cloud Technologies: GCP (BigQuery, GCS, Dataproc, Vertex AI); AWS (S3, Redshift), Snowflake
ETL and Data Pipelines: Apache Airflow, dbt, Kafka, Informatica, Google Cloud Composer, Git, Azure DevOps, CI/CD
AI & Analytics: Power BI, Tableau, ThoughtSpot, Looker, XGBoost; Terraform; LangChain, embeddings, RAG, Root cause analysis